

de.NBI Cloud Usermeeting - 2023

Introduction to Kubernetes I: Basic Concepts

Sebastian Beyvers
sebastian.beyvers@cb.jlug.de

27 November 2023



What is Kubernetes (K8s) ?



kubernetes

- Container orchestration framework
- Open sourced by Google in 2014: Managed by CNCF
 - Based on a Google internal framework called Borg
- The de-facto standard container orchestration framework
 - Adopted by most cloud providers: GCP (GKE), AWS (EKS), Azure (AKS),...
- Other competing frameworks got pushed out:
 - Apache Mesos -> mainly for deploying infrastructure
 - Docker swarm -> nowadays irrelevant

Kubernetes architecture

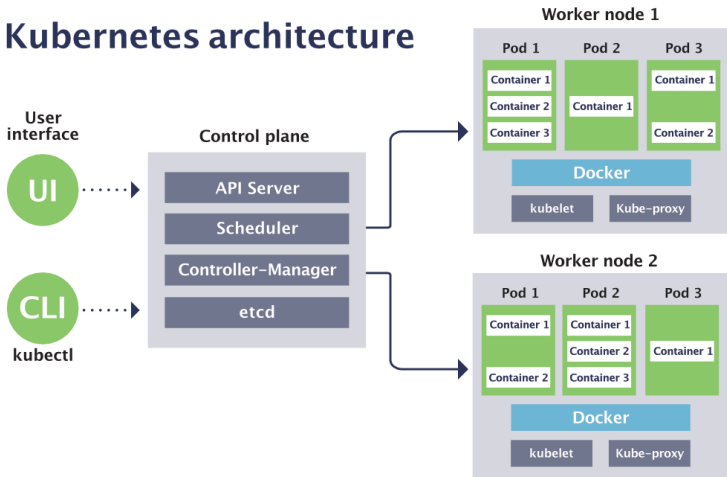


Figure 1: Basic overview of the infrastructure layout of Kubernetes Source: <https://sensu.io/blog/how-kubernetes-works>

Kubernetes - Resources

- Everything in Kubernetes is a resource
- Resources are managed by dedicated controllers
- You can extend K8s with own resources and controllers (CRDs)

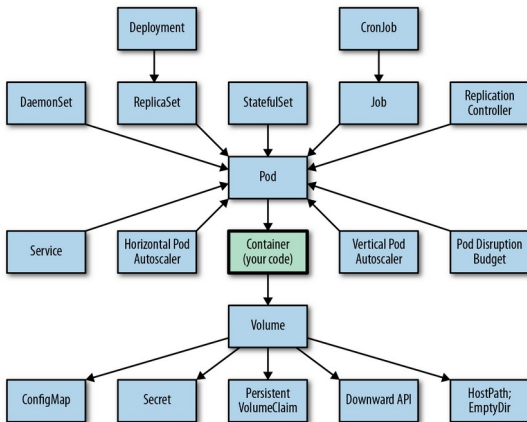


Figure 2: Basic overview of resources that interact with the pod resource:
<https://dev4devs.com/2019/10/20/what-are-the-kubernetes-resources-which-are-most-useful-for-developers/>

Kubernetes – Core Concepts I

- **Infrastructure**

- ControlPlane-Worker (ControlPlane – Node) architecture
 - ControlPlane runs API, etcd, Scheduler, ...
 - Nodes run kubelet, kube-proxy, ...

- **Deployment units**

- Pod
 - Basic deployment unit
 - Runs 1..n container
- Replica set
 - Creates 0..n replicas of Pods
- Deployment
 - Manages replica sets
 - Helps with upgrading (rolling updates/rollbacks)
- Stateful set
 - Pods that stay on a node and retain "state" -> Databases
- Jobs & Cronjobs
 - Pods that run to "completion" -> Workflows
- ...

Kubernetes - Concepts II

- **Networking**

- Flat overlay network
 - By default: Every pod can see every other pod
- Each pod has an individual (non-public) IP

- **High-level organization**

- Service: pods can be grouped by a single name -> (internal) DNS
- Ingress: Routes external traffic
 - Multiple implementation, in our case: nginx ingress

- **Scheduling**

- Fully automated based on resource requests
- Resources:
 - CPU -> 0.5 CPU is a valid resource requests/quota
 - RAM
 - Storage -> Can separate various storage classes
 - GPUs

Kubernetes - Concepts III

- RBAC for tenancy and authorization
- Namespaces can be used to support tenancy
- Most resources created by users are bound to a namespace
 - Some (especially internal) resources are not bound to a namespace
- Namespaces do not handle multi-tenancy well
 - Addons and distros can add additional capabilities
- Containers (Docker) also have drawbacks for hard multi-tenancy
 - Docker got mostly replaced by plain containerd

- **RESTful API**

- Uses standard HTTP verbs to define actions (POST, DELETE, UPDATE, ...)

- **Verbs applied to resource types (pod, deployments, ...)**

- Resource types are represented as “kind”
- kinds are strictly versioned (mostly semantic versioning)

- **Kubectl: CLI**

- Verbs are slightly different (create, get, describe, delete, ...)
- Similar syntax -> kubectl verb resource [options]
- Example:
 - kubectl get pods -n namespace
- Create command often used along with YAML files
 - YAML file contain declarative resource configurations

Excursus - YAML

YAML Ain't Markup Language is a human-readable data-serialization language and super-set of JSON.

- Indentation matters (similar to Python)
- You can separate resources with three dashes ---
- Hierarchical structure, - indicates lists.

```
apiVersion: batch/v1
kind: Job
metadata:
  name: hello
spec:
  template:
    # This is the pod template
    spec:
      containers:
        - name: hello
          image: busybox
          command: ['sh', '-c', 'echo "Hello, Kubernetes!" && sleep 3600']
      restartPolicy: OnFailure
    # The pod template ends here
```

Pods




Figure 4: Basic overview of a Job. Source: <https://kubernetes.io/docs/tutorials/kubernetes-basics/explore/explore-intro/>

Additional features for pods:

- Liveness probes
- Init containers
- Lifecycle: Pending, Running, Succeeded, Failed, Unknown

Jobs and Cronjobs

Jobs are pods that run to completion until the status is either *succeeded* or *failed*. **CronJobs** are jobs that run in a periodic interval.

```
apiVersion: batch/v1
kind: CronJob
metadata:
  name: hello
spec:
  schedule: "*/1 * * * *"
  jobTemplate:
    spec:
      template:
        spec:
          containers:
            - name: hello
              image: busybox
              imagePullPolicy: IfNotPresent
              command:
                - /bin/sh
                - -c
                - date; echo Hello from the Kubernetes cluster
          restartPolicy: OnFailure
```

Excursus - Kubermatic WebUI

The screenshot displays the Kubermatic WebUI interface. At the top, there's a navigation bar with the Kubermatic logo, 'KUBERMATIC Kubernetes Platform', and tabs for 'Projects' and 'Schulungen'. On the right, there are icons for a globe, a bell, and a user profile. A left sidebar contains icons for a grid, a cluster, a calendar, and a key. The main content area shows the details for a cluster named 'kurs-2023'. A red box highlights the 'Get Kubeconfig' button. Below this, there's a table of cluster details:

Control Plane	CNI Plugin	Cluster ID	Created	Seed	Region
1.26.9	canal v3.24	hw5hcs4298	a year ago	kubermatic	DE (Giessen)

Below the table, there are sections for 'Provider' (openstack), 'Container Runtime' (containerd), 'SSH Keys' (No assigned keys), 'Nodes CPU Usage' (454/68000 millicores), 'Nodes Memory Usage' (6218/157990 MiB), 'Control Plane CPU Usage' (685 millicores), and 'Control Plane Memory Usage' (2576 MiB). At the bottom, there's a link to 'ADDITIONAL CLUSTER INFORMATION'.

- Kubermatic is one of many Kubernetes distributions
 - Easier cluster deployment
 - Web GUI for Users and Admins
 - Openstack integration to automatically provision clusters
- For now we use it to get the `kubectl` config
- <https://www.k8s.gi.denbi.de>

Questions?